

# 64-bit address spaces

- **Straight hierarchical page tables not efficient**
- **Solution 1: Guarded page tables [Liedtke]**
  - Omit intermediary tables with only one entry
  - Add predicate in high level tables, stating the only virtual address range mapped underneath + # bits to skip
- **Solution 2: Hashed page tables**
  - Store Virtual  $\rightarrow$  Physical translations in hash table
  - Table size proportional to physical memory
  - Clustering makes this more efficient

# OS policy choices (besides page table)

- **Page replacement**

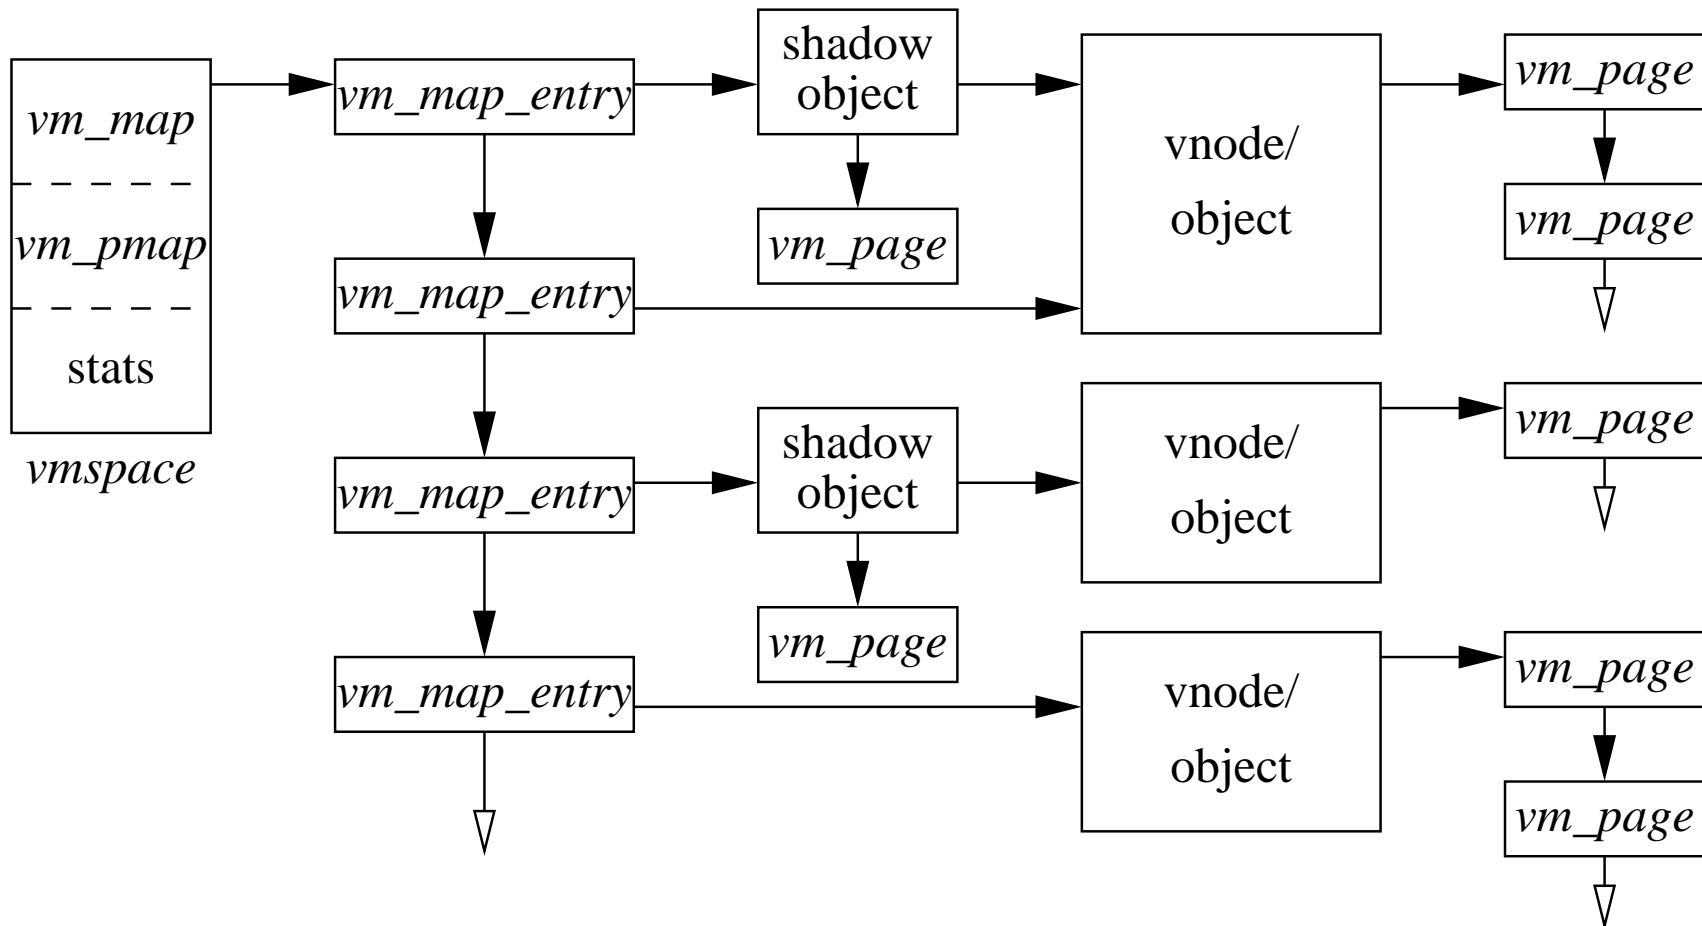
- Optimal – Least soon to be used (impossible)
- Least recently used (hard to implement)
- Random
- Not recently used

- **Direct-mapped physical caches**

- Virtual  $\rightarrow$  Physical mapping can affect performance
- Applications can conflict with each other or themselves
- Scientific applications benefit if consecutive virtual pages to not conflict in the cache
- Many other applications do better with random mapping

## 4.4 BSD VM system

- **Each process has a *vm\_space* structure containing**
  - *vm\_map* – machine-independent virtual address space
  - *vm\_pmap* – machine-dependent data structures
  - statistics – e.g. for syscalls like *getrusage()*
- ***vm\_map* is a linked list of *vm\_map\_entry* structs**
  - *vm\_map\_entry* covers contiguous virtual memory
  - points to *vm\_object* struct
- ***vm\_object* is source of data**
  - e.g. vnode object for memory mapped file
  - points to list of *vm\_page* structs (one per mapped page)
  - *shadow objects* point to other objects for copy on write



# Pmap (machine-dependent) layer

- **Pmap layer holds architecture-specific VM code**
- **VM layer invokes pmap layer**
  - On page faults to install mappings
  - To protect or unmap pages
  - To ask for dirty/accessed bits
- **Pmap layer is lazy and can discard mappings**
  - No need to notify VM layer
  - Process will fault and VM layer must reinstall mapping
- **Pmap handles restrictions imposed by cache**

# Example uses

- ***vm\_map\_entry* structs for a process**
  - r/o text segment → file object
  - r/w data segment → shadow object → file object
  - r/w stack → anonymous object
- **New *vm\_map\_entry* objects after a fork:**
  - Share text segment directly (read-only)
  - Share data through two new shadow objects  
(must share pre-fork but not post fork changes)
  - Share stack through two new shadow objects
- **Must discard/collapse superfluous shadows**
  - E.g., when child process exits

# What happens on a fault?

- **Traverse *vm\_map\_entry* list to get appropriate entry**
  - No entry? Protection violation? Send process a SIGSEGV
- **Traverse list of [shadow] objects**
- **For each object, traverse *vm\_page* structs**
- **Found a *vm\_page* for this object?**
  - If first *vm\_object* in chain, map page
  - If read fault, install page read only
  - Else if write fault, install copy of page
- **Else get page from object**
  - Page in from file, zero-fill new page, etc.

# DEC Alpha MMU

- **Software managed TLB (like MIPS)**
  - In model used by paper: 8K base pages, TLB supports 128 instruction/128 data entries
- **But TLB miss handler not part of OS**
  - Processor ships with special “PAL code” in ROM
  - Processor-specific, but provides uniform interface to OS
  - Paper calls this firmware, but runs from memory like OS
- **Various events vector directly to PAL code**
  - CALL\_PAL instruction, TLB miss/fault, FP disabled
- **PAL code runs in special privileged processor mode**
  - Interrupts always disabled
  - Have access to special instructions and registers



# PAL code interface details

- **Examples of Digital Unix PALcode entry functions**
  - `callsys/retsys` - make, return from system call
  - `swpctx` - change address spaces
  - `wrvptptr` - write virtual page table pointer
  - `tbi` - TBL invalidate
- **Some fields in PALcode page table entries**
  - GH - 2-bit granularity hint  $\rightarrow 2^N$  pages have same translation
  - ASM - address space match  $\rightarrow$  mapping applies in all processes